

Just lg

June Lemieux

2024-12-13

```
# libraries/etc  
library(tidymodels)
```

```
## -- Attaching packages ----- tidymodels 1.2.0 --
```

```
## v broom      1.0.7      v recipes      1.1.0  
## v dials      1.3.0      v rsample      1.2.1  
## v dplyr      1.1.4      v tibble      3.2.1  
## v ggplot2    3.5.1      v tidyr      1.3.1  
## v infer      1.0.7      v tune        1.2.1  
## v modeldata  1.4.0      v workflows   1.1.4  
## v parsnip    1.2.1      v workflowsets 1.1.0  
## v purrr      1.0.2      v yardstick   1.3.1
```

```
## -- Conflicts ----- tidymodels_conflicts() --
```

```
## x purrr::discard() masks scales::discard()  
## x dplyr::filter()  masks stats::filter()  
## x dplyr::lag()     masks stats::lag()  
## x recipes::step() masks stats::step()  
## * Use suppressPackageStartupMessages() to eliminate package startup messages
```

```
library(ggplot2)  
library(vip)
```

```
## Warning: package 'vip' was built under R version 4.4.2
```

```
##
```

```
## Attaching package: 'vip'
```

```
## The following object is masked from 'package:utils':
```

```
##
```

```
##      vi
```

```
options(scipen=4)
```

```
## Data Processing
```

```
# Initial data import
```

```
customer_data <- read.csv("~/00 Merrimack/DSE6111/Course Project/customer_data_1.csv", header = TRUE)
```

```
customer_data$Attrition_Flag <- ifelse(customer_data$Attrition_Flag == "Attrited Customer", 1, 0)
customer_data$Attrition_Flag <- factor(customer_data$Attrition_Flag)
```

```
# check levels
levels(customer_data$Attrition_Flag)
```

```
## [1] "0" "1"
```

```
# set seed
set.seed(10)
```

```
# get rid of ClientNum identifier
customer_data <- customer_data[ -c(1) ]
```

```
# create data split
data_split <- initial_split(customer_data, strata = "Attrition_Flag", prop = 0.75)
```

```
# create training and test sets
training_set <- training(data_split)
test_set <- testing(data_split)
```

```
# row count per set
nrow(training_set)
```

```
## [1] 7595
```

```
nrow(test_set)
```

```
## [1] 2532
```

```
# use cv folds and tune
cv_set <- vfold_cv(training_set, v = 10)
```

```
# create recipe for preprocessing - only numeric var at first
attrition_recipe_lg <-
  recipe(
    Attrition_Flag ~ Months_on_book + Total_Relationship_Count +
      Months_Inactive_12_mon + Contacts_Count_12_mon + Total_Amt_Chng_Q4_Q1 +
      Total_Trans_Amt + Total_Trans_Ct + Total_Ct_Chng_Q4_Q1 +
      Avg_Utilization_Ratio,
    data = training_set,
  ) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_center(all_predictors()) %>%
  step_scale(all_predictors())
```

```
# model with lg
lg_spec <-
  logistic_reg(penalty = tune(), mixture = 1) %>%
  set_engine("glmnet") %>%
```

```

set_mode("classification")

# create workflow
lg_workflow <- workflow() %>%
  add_recipe(attrition_recipe_lg) %>%
  add_model(lg_spec)

# build tuning grid
lg_tune_grid <- tibble(penalty = 10^seq(-1, 0, length.out = 15))

lg_results <-
  lg_workflow %>%
  tune_grid(
    cv_set,
    grid = lg_tune_grid,
    control = control_grid(save_pred = TRUE),
    metrics = metric_set(roc_auc, accuracy)
  )

```


```
## Warning: package 'glmnet' was built under R version 4.4.2
```

```
lg_results %>% show_best(metric = "roc_auc")
```

```
## # A tibble: 5 x 7
```

	penalty	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
## 1	0.1	roc_auc	binary	0.803	10	0.00420	Preprocessor1_Model01
## 2	0.118	roc_auc	binary	0.796	10	0.00429	Preprocessor1_Model02
## 3	0.139	roc_auc	binary	0.5	10	0	Preprocessor1_Model03
## 4	0.164	roc_auc	binary	0.5	10	0	Preprocessor1_Model04
## 5	0.193	roc_auc	binary	0.5	10	0	Preprocessor1_Model05


```
autoplot(lg_results)
```



```
best_auc <- lg_results %>% select_best(metric = "roc_auc")
best_acc <- lg_results %>% select_best(metric = "accuracy")
```

```
lg_plot <-
  lg_results %>%
  collect_metrics() %>%
  ggplot(aes(x = penalty, y = mean)) +
    geom_point() +
    geom_line() +
    ylab("Area under the ROC curve") +
    scale_x_log10(labels = scales::label_number()) +
    ggtitle("Plot of lg_results", subtitle="AUC")
```

```
lg_plot
```




```
final_lg_workflow <- finalize_workflow(
  lg_workflow,
  best_auc
)

last_lg_fit <-
  final_lg_workflow %>%
  last_fit(data_split)

lg_metrics <- last_lg_fit %>% collect_metrics()
```

```
# feature importance -
last_lg_fit %>%
  extract_fit_parsnip() %>%
  vip(num_features = 10)
```



```
# compare to the null model
null_classification <- null_model() %>%
  set_engine("parsnip") %>%
  set_mode("classification")

null_lg <- workflow() %>%
  add_recipe(attrition_recipe_lg) %>%
  add_model(null_classification) %>%
  fit_resamples(
    cv_set
  )

null_metrics <- null_lg %>% collect_metrics()

##### plot ROC curve #####
test_results <-
  test_set %>%
  select(Attrition_Flag) %>%
  bind_cols (
    last_lg_fit %>% collect_predictions() %>%
    select(p_1 = .pred_1)
  )

roc_data <- data.frame(threshold = seq(1, 0, -0.01), fpr = 0, trp = 0)
for (i in roc_data$threshold) {
  over_threshold <- test_results[test_results$p_1 >= i, ]
}
```


```

fpr <- sum(over_threshold$Attrition_Flag == 0)/sum(test_results$Attrition_Flag==0)
roc_data[roc_data$threshold==i, "fpr"] <- fpr

tpr <- sum(over_threshold$Attrition_Flag == 1)/sum(test_results$Attrition_Flag==1)
roc_data[roc_data$threshold==i, "tpr"] <- tpr
}

ggplot() +
  geom_line(data = roc_data, aes(x = fpr, y = tpr, color = threshold), linewidth = 2) +
  scale_color_gradientn(colors = rainbow(3)) +
  geom_abline(intercept = 0, slope = 1, lty = 2) +
  geom_point(data = roc_data[seq(1, 101, 10), ], aes(x = fpr, y = tpr)) +
  geom_text(data = roc_data[seq(1, 101, 10), ],
            aes(x = fpr, y = tpr, label = threshold, hjust = 1.2, vjust = -0.2)) +
  ggtitle("ROC Curve", subtitle = "Lasso Logistic")

```




```

##### ROC curve calculation breakdown #####

ggplot(data = test_results, aes(x = p_1, y = Attrition_Flag)) +
  geom_jitter()

```



```
threshold <- 0.23


test_results$predictions <- ifelse(test_results$p_1 >= threshold, 1, 0)
tp <- nrow(test_results[test_results$Attrition_Flag==1 & test_results$predictions==1, ])
fp <- nrow(test_results[test_results$Attrition_Flag==0 & test_results$predictions==1, ])
tn <- nrow(test_results[test_results$Attrition_Flag==0 & test_results$predictions==0, ])
fn <- nrow(test_results[test_results$Attrition_Flag==1 & test_results$predictions==0, ])

test_results$type <- ""
test_results[test_results$Attrition_Flag==1 & test_results$predictions==1, "type"] <- "tp"
test_results[test_results$Attrition_Flag==0 & test_results$predictions==1, "type"] <- "fp"
test_results[test_results$Attrition_Flag==0 & test_results$predictions==0, "type"] <- "tn"
test_results[test_results$Attrition_Flag==1 & test_results$predictions==0, "type"] <- "fn"

ggplot(data = test_results, aes(x = p_1, y = Attrition_Flag)) +
  geom_jitter(aes(colour = type)) +
  geom_vline(xintercept = threshold, linetype = "dashed", color = "blue", linewidth = 1.5) +
  scale_color_brewer(palette = "RdYlBu") +
  ggtitle("ROC Curve Calculation Breakdown", subtitle = "Lasso Logistic")
```


ROC Curve Calculation Breakdown

Lasso Logistic



```
fpr <- fp/(fp + tn)
tpr <- tp/(tp + fn)
```

```
# confusion matrix data
tn
```

```
## [1] 2104
```

```
fn
```

```
## [1] 381
```

```
fp
```

```
## [1] 21
```

```
tp
```

```
## [1] 26
```

```
““
```